



MASTER OF SCIENCE
IN ENGINEERING

Machine Learning

T-MachLe

7. Support Vector Machines (SVM)

Christophe Gisler
Jean Hennebert
Andres Perez Uribe

Plan - Classification task with Support Vector Machines (SVM)


- 7.1 Recaps on classification task and logistic regression
- 7.2 Linear SVM for linearly separable data
- 7.3 Linear SVM for not linearly separable data
- 7.4 Nonlinear SVM for not linearly separable data
- 7.5 Multiclass SVM
- 7.6 History & References

Practical Work 7

7.1 Recaps On classification tasks and logistic regression


Classification task

- Example:
 - **2 classes:**
 - Women (●)
 - Men (●)
 - **2 features:**
 - Hair length (axe X)
 - Weight (axe Y)
 - **What about the new sample (●)?**
Is it a woman or a man?




Classification task - generative models

- Some algorithms try to build a **model per class**, i.e. generative model explaining the distribution within the class
 - E.g. Likelihood estimation with Gaussian Mixture Models (**GMM**)




Classification task - discriminant models


- Some algorithms try to **discriminate classes**, i.e. to build a border defining the classes
- Last week we have seen **Logistic Regression**
 - Discover the equation of the border that maximise the correct classifications
 - The equation of the border is feeding a sigmoid so that on one side of the border the probability value is larger than 0.5 (class 1) or smaller (class 0)
 - Logistic regression may find many possibilities of hyperplanes
- Which is the **best separating hyperplane**?



$$\hat{y} = P(C_1 | x) = 1 - P(C_2 | x)$$
$$\hat{y} = h_{\theta}(\mathbf{x}) = \text{sigmoid}(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$




Logistic regression with non-linear decision boundaries




$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

(g = sigmoid function)



$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2 + \theta_5 x_1 x_2)$$



$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_1^2 + \theta_3 x_1^2 x_2 + \theta_4 x_1^2 x_2^2 + \theta_5 x_1^2 x_2^3 + \theta_6 x_1^3 x_2 + \dots)$$

- To move to non-linear decision boundaries, just add features to the x array
 - e.g. representing dependencies to the squared

$$h_{\theta}(\mathbf{x}) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \dots)$$



$$h_{\theta}(\mathbf{x}) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \dots)$$

with $x_3 = x_1^2$

x_1 = surf	x_2 = rooms	x_3 = surf ²	y = rented
26	1	26*26 =	0
37	1.5	1369	1
57	2	3249	0
48	2	2304	1
...

Training discriminant models

- A common strategy to find the best parameters through training is the **gradient descent**.

Gradient descent involves the computation of two mathematical terms:

- A **loss function** $J(\theta)$
 - It expresses how bad we are with the current values of parameters θ on the train set
 - We want to minimise this function through the training procedure
- The **gradient of the loss** with respect to the parameters $\frac{\partial J}{\partial \theta}$

- The negative of the gradient will tell us in which direction to move the parameters to minimise the loss

Loss function and gradient for logistic regression

- We want to maximise the number of correct classifications, i.e. the product over the training set of the a posteriori probabilities.
 - As the derivation of the product is not easy, we took the log of the product which transformed as a sum of the logs.

$$J(\theta) = -\frac{1}{N} \sum_{n=1}^N J(\theta, x_n) = -\frac{1}{N} \sum_{n=1}^N y_n \log h_{\theta}(\mathbf{x}_n) + (1 - y_n) \log(1 - h_{\theta}(\mathbf{x}_n))$$

Note : in the last class, the negative term was not there and we had to maximise the function J

When $y_n = 1$ $J(\theta, x_n) = -\log(h_{\theta}(\mathbf{x}_n)) = -\log(\hat{y}_n)$

When $y_n = 0$ $J(\theta, x_n) = -\log(1 - h_{\theta}(\mathbf{x}_n)) = -\log(1 - \hat{y}_n)$

- After some mathematical developments, the gradient of the loss w.r.t. the parameter is

$$\frac{\partial J(\theta)}{\partial \theta_i} = \frac{1}{N} \sum_{n=1}^N (y_n - h_{\theta}(\mathbf{x}_n)) x_{n,i}$$

Target
value

Gotten
output

Training with a gradient descent

- Update any parameters of your model in the opposite direction of the gradient of the loss w.r.t. weights

$$\text{param} \leftarrow \text{param} - \alpha \frac{\partial J}{\partial \text{param}}$$

In practice we have stopping criteria, e.g. `epoch < max_epochs` or `loss_gain < ε`

1 sample = stochastic gradient descent
B samples = mini-batch gradient descent
N samples = (full) batch gradient descent

Current weights

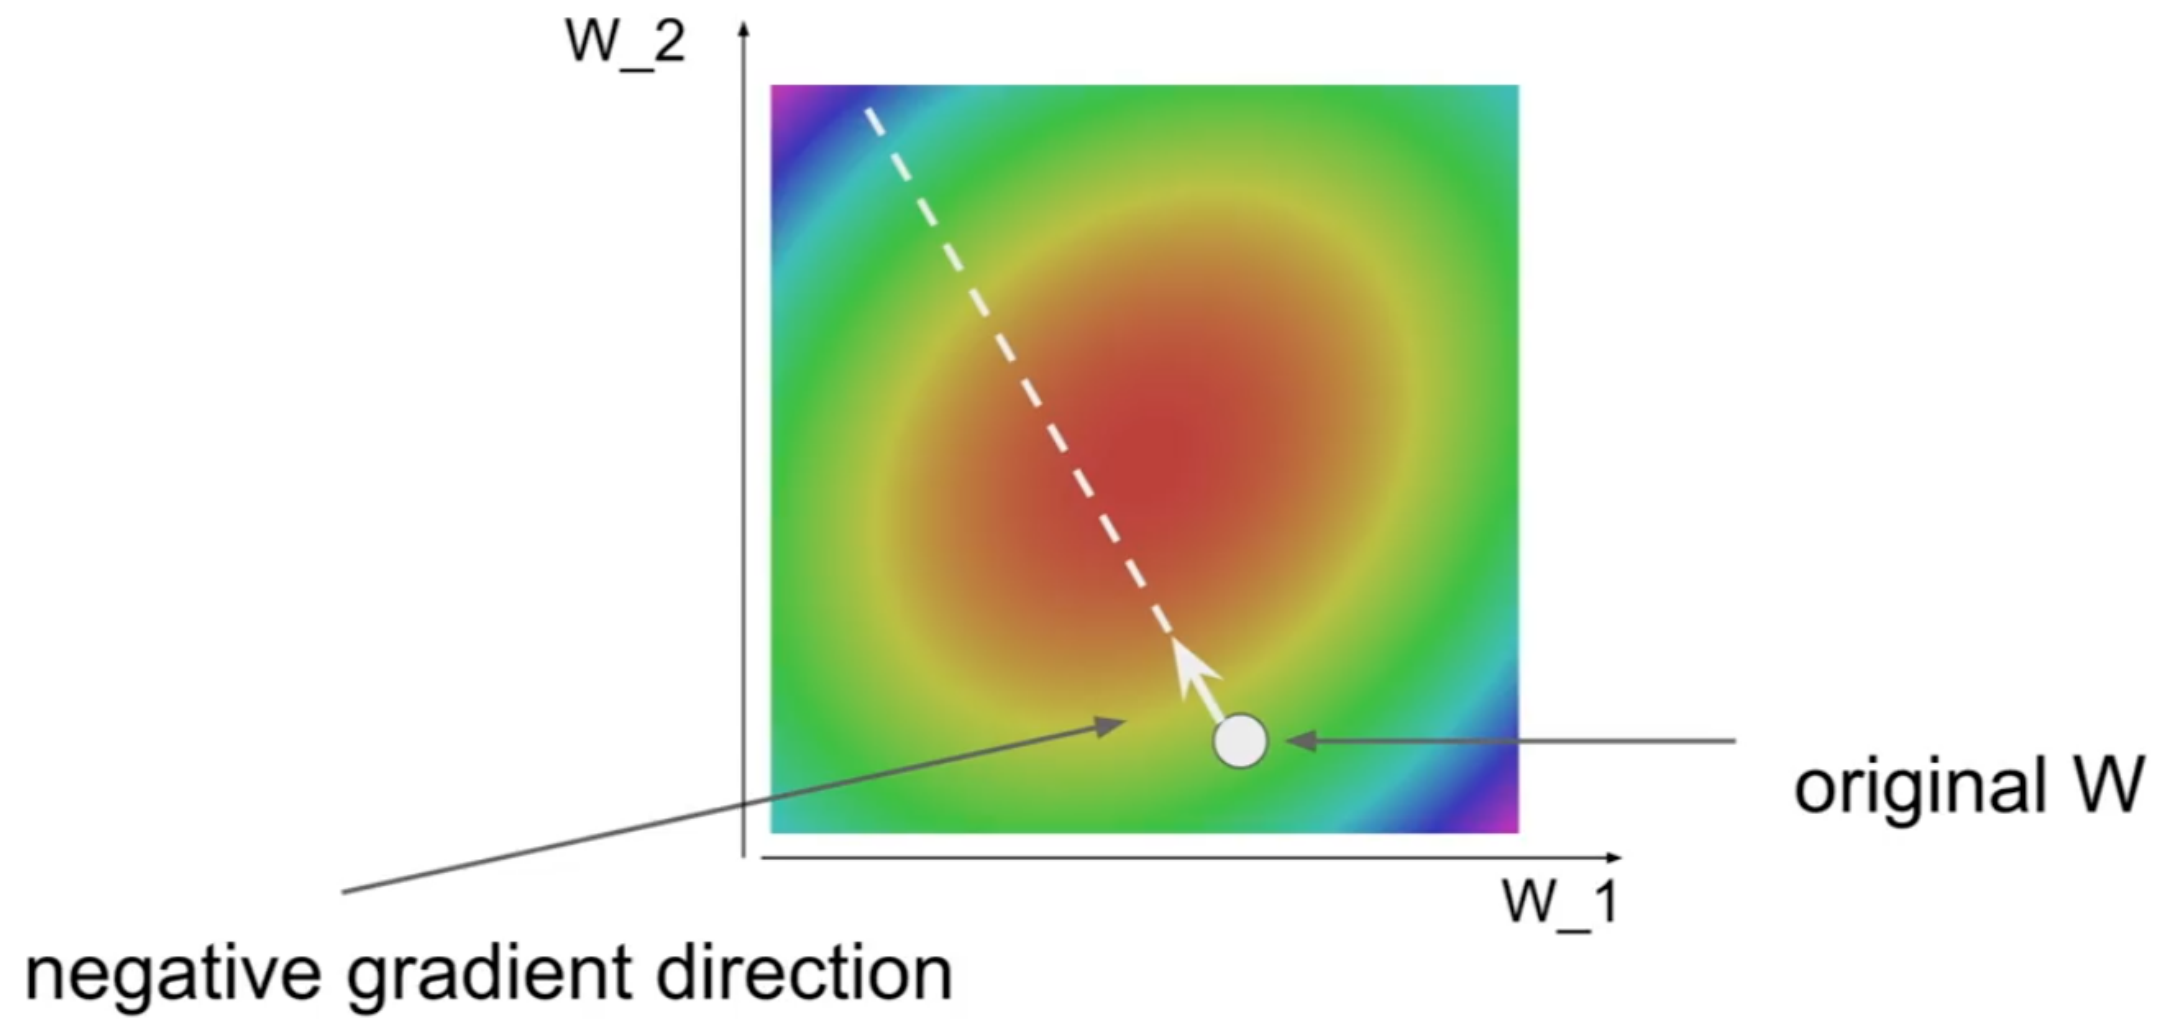
```
# Vanilla Gradient Descent
```

```
while True:
```

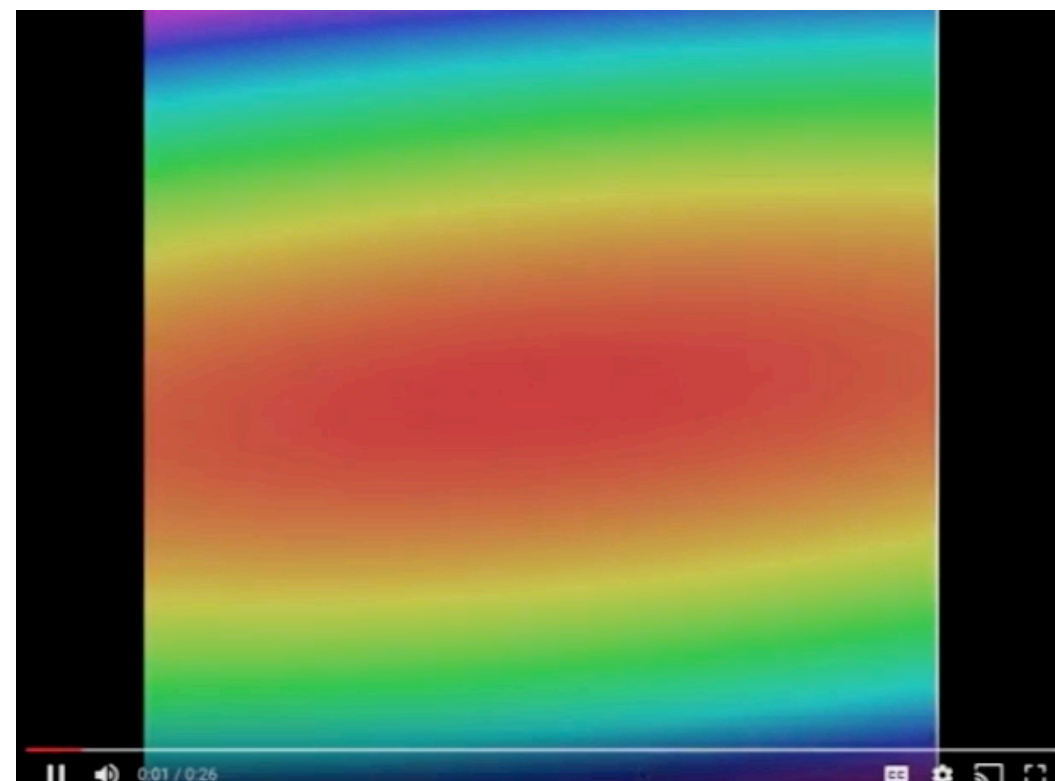
```
    weights_grad = evaluate_gradient(loss_fun, data, weights)
```

```
    weights += - step_size * weights_grad # perform parameter update
```

The loss function



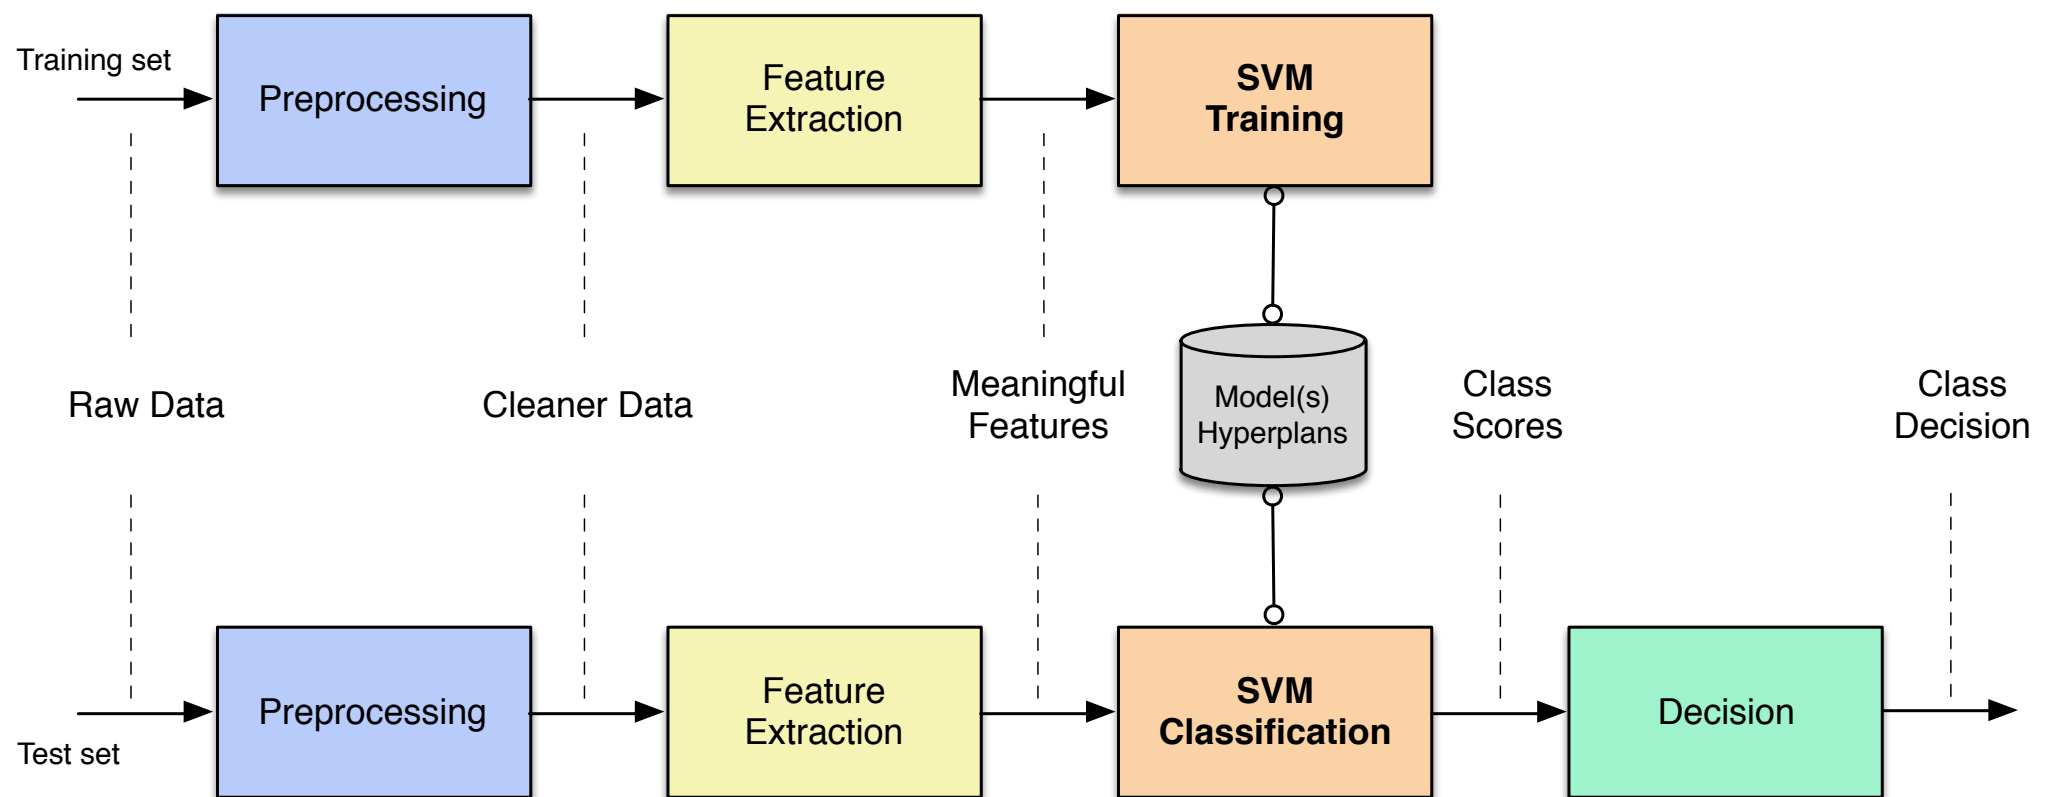
From Fei-Fei Li & Justin Johnson & Serena Yeung, April 2018: CS231n Stanford



7.2 Linear SVM

Classification task

TRAINING



TEST

Linear SVM

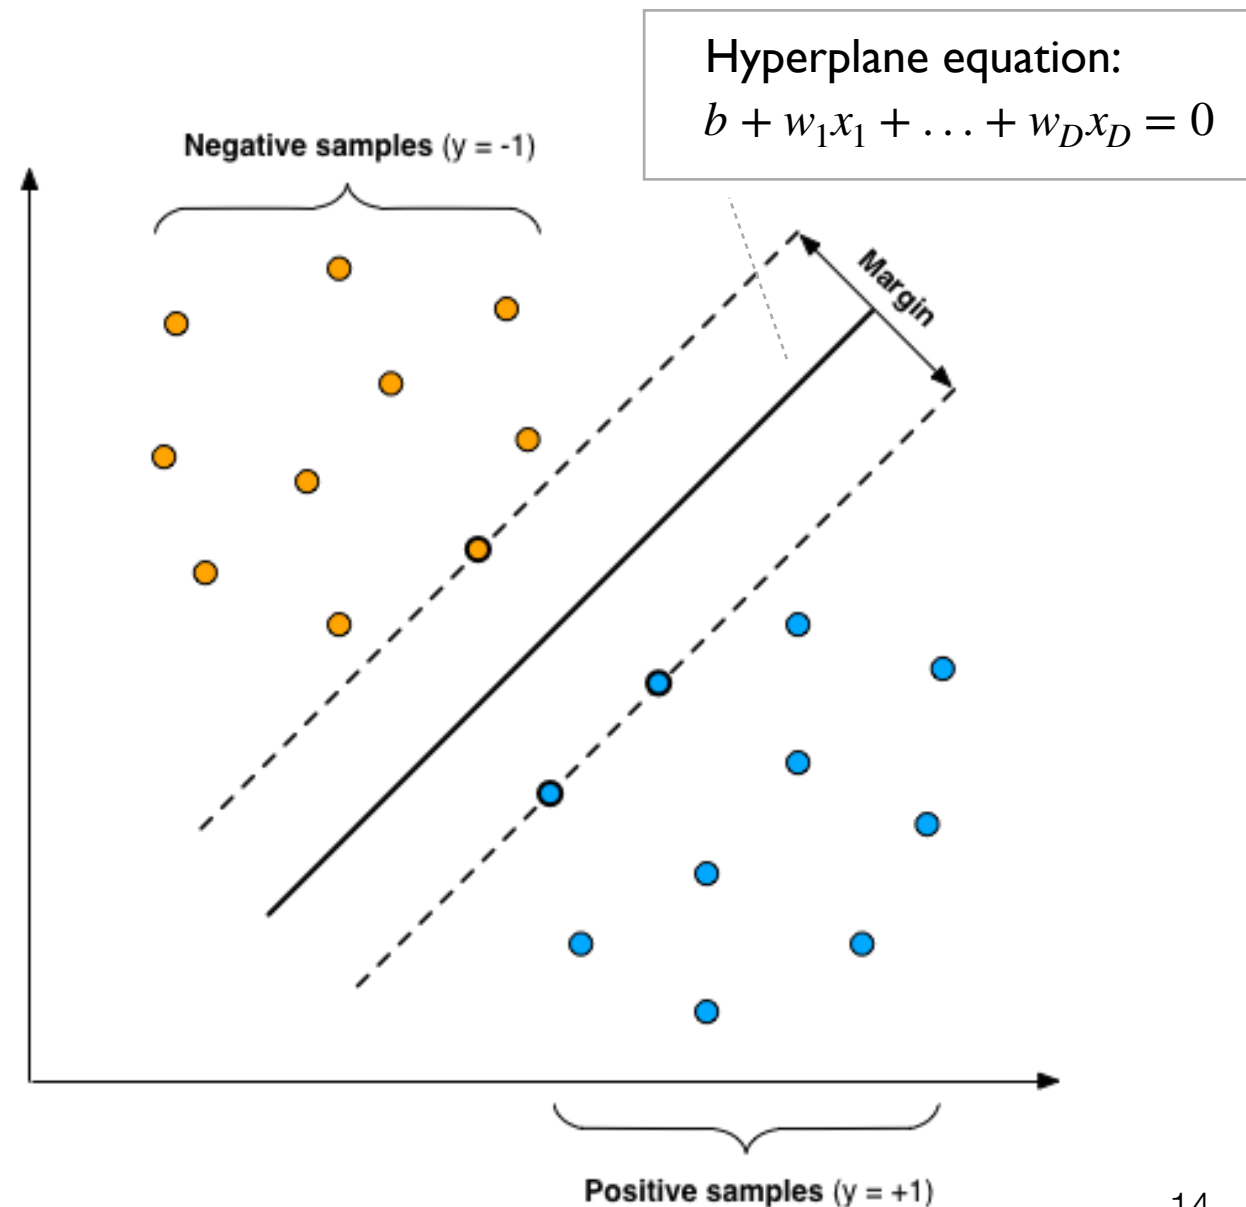

- Given:
 - Training samples $\mathbf{x}_i \in \mathbb{R}^D$
 - A two class problems with targets $y_i \in \{-1, +1\}$ associated to classes C_1, C_0
 - Number of samples : N with $i = 1, \dots, N$
 - The hypothesis function:

$$h_w(\mathbf{x}) = \text{sign}(b + w_1x_1 + \dots + w_Dx_D)$$

$$h_w(\mathbf{x}) = \text{sign}(b + \mathbf{w}\mathbf{x})$$

A linear SVM tries to find the **hyperplane** that **separates** the 2 classes and that **maximizes** the **margin** between the 2 classes

In previous classes, we used symbol θ instead of \mathbf{w}



Linear SVM


What does it mean to maximise the margin?

- Margin = distance M between the 2 parallel hyperplanes (on boundaries):
 $\mathbf{w} \cdot \mathbf{x} + b = -1$ (a is a constant >0)
 $\mathbf{w} \cdot \mathbf{x} + b = +1$
- Let's define $\mathbf{x}_1, \mathbf{x}_2$ as points on these 2 hyperplanes such that: $\mathbf{x}_1 - \mathbf{x}_2 = t \cdot \mathbf{w}$ (t is a scalar)
- So we have: $\mathbf{x}_2 = \mathbf{x}_1 + t \cdot \mathbf{w}$
 $M = \|\mathbf{x}_1 - \mathbf{x}_2\| = \|t \cdot \mathbf{w}\| = t \cdot \|\mathbf{w}\|$ (1)
 $\mathbf{w} \cdot \mathbf{x}_1 + b = -1$ (2)
 $\mathbf{w} \cdot \mathbf{x}_2 + b = +1$ (3)
- (3) - (2) gives: $\mathbf{w} \cdot (\mathbf{x}_2 - \mathbf{x}_1) = 2$
 $\Leftrightarrow \mathbf{w} \cdot (\mathbf{x}_1 + t \cdot \mathbf{w} - \mathbf{x}_1) = 2$
 $\Leftrightarrow \mathbf{w} \cdot t \cdot \mathbf{w} = 2 \Leftrightarrow t (\mathbf{w} \cdot \mathbf{w}) = 2$
 $\Leftrightarrow t \|\mathbf{w}\|^2 = 2 \Leftrightarrow t = 2 / \|\mathbf{w}\|^2$ (4)
- (4) in (1) gives: $M = 2 / \|\mathbf{w}\|$
- Hence to maximize M , we have to minimize $\|\mathbf{w}\|$ while respecting the correct classification constraints (see next slide)

By geometric definition \mathbf{w} is the normal vector of the hyperplane, i.e. perpendicular to the hyperplane

$\|\mathbf{x}\|$ is the Euclidian norm of vector \mathbf{x}


$$\|\mathbf{x}\|_2 := \sqrt{x_1^2 + \dots + x_n^2}.$$



In other words, a term on the loss function will include $\|\mathbf{w}\|$


Linear SVM

- Training samples on the margin boundaries are called the **support vectors**



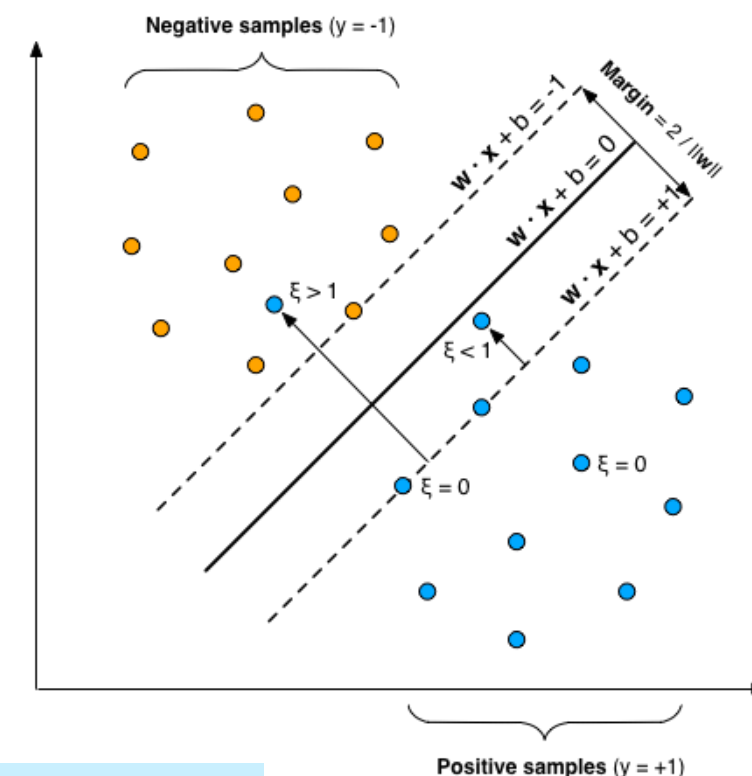

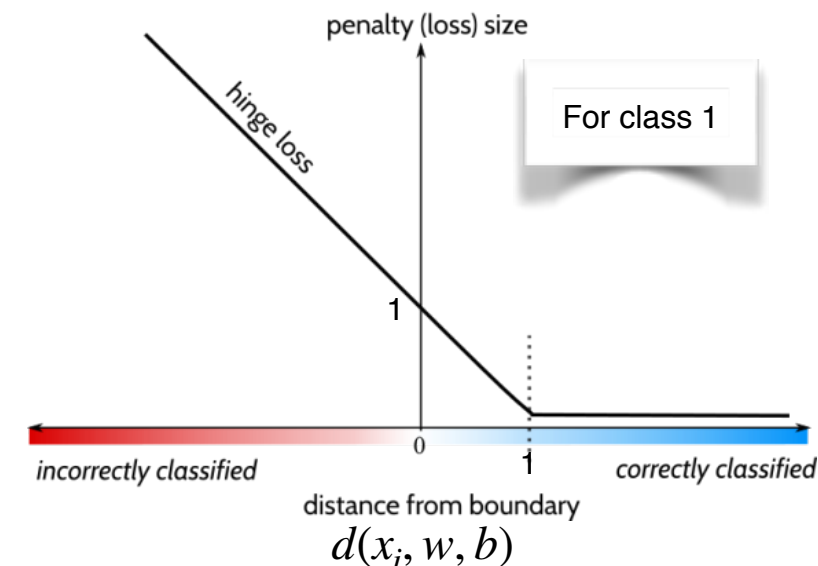
Linear SVM for not linearly separable data

- **Problem:** how to linearly separate these 2 classes?
- We need to inject another term in the cost function that will minimise the number of incorrectly classified samples.
- For SVM, the “Hinge” loss is used
 - It takes into account the notion of margin



Linear SVM for not linearly separable data

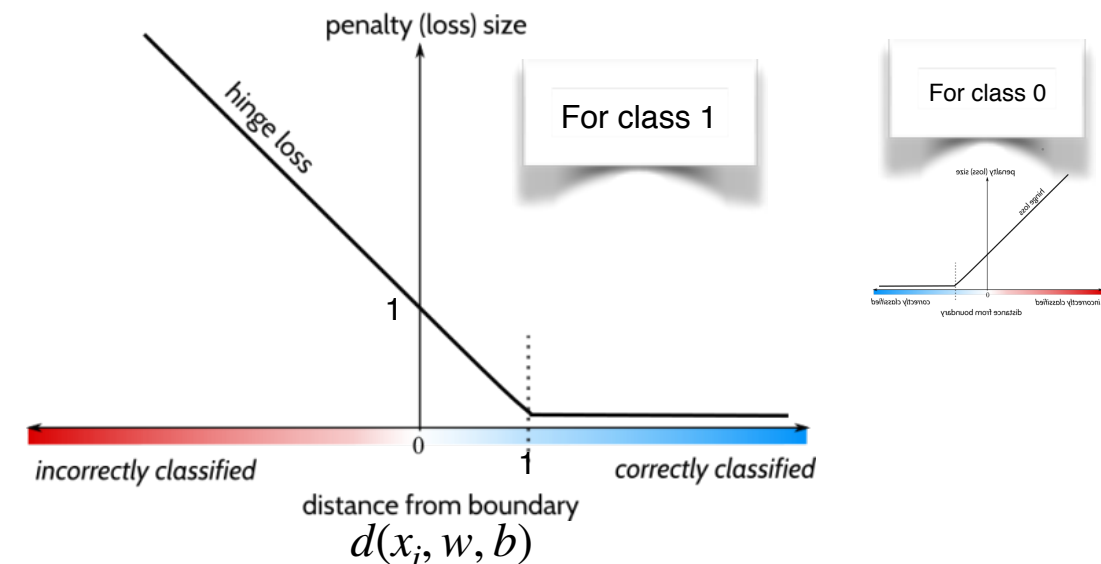
- For each sample \mathbf{x}_i , compute the Hinge loss ξ_i
 - = 0 if the point falls above the margin
 - = 1 if the point falls on the hyperplane
 - > 1 if the point is on the wrong side of the hyperplane
- The $\xi \geq 0$ measure the degree of misclassification of samples in terms of distance to the plane
- Overall loss is quantified by: $\sum \xi_i$ for $i=1, \dots, N$
- In the SVM terminology, the ξ_i are also called the **slack** variables.



In other words, a term on the loss function will include $\sum \xi_i$

SVM loss function

- We can then define the SVM loss function taking into account of two terms to minimise:
 - One based on the Hinge loss for both classes C_1 and C_0
 - One based on $\|\mathbf{w}\|$
 - For mathematical reasons, we instead minimise $\|\mathbf{w}\|^2/2$



$$J(\theta) = C \left[\frac{1}{N} \sum_{i=1}^N y_i \text{hinge}_{C_1}(d(x_i, w, b)) + (1 - y_i) \text{hinge}_{C_0}(d(x_i, w, b)) \right] + \left[\frac{1}{2} \|\mathbf{w}\|^2 \right]$$

Trade-off coef : giving more or less importance to perf term

Performance term: on the data, how far are we predicting from the ground truth?

Regularisation term: impeach too large values of \mathbf{w}

Impact of parameter C

- The factor C is a regularization parameter which trades off the margin size and the training error
- The smaller C , the greater the number of admitted misclassified train samples

 $C=100$  $C=1$  $C=0.15$  $C=0.1$



Minimizing the loss function

- 2 possibilities
- Use math toolboxes for minimisation problems under constraints
 - For example: SciKit Learn **`svm.SVC()`** based on the popular library libsvm
 - <http://scikit-learn.org/stable/modules/svm.html>
 - <http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html#sklearn.svm.SVC>
 - <https://www.csie.ntu.edu.tw/~cjlin/libsvm/>
 - Practical considerations: usually less tuning to perform, libsvm is well optimised and stabilised library
- Use gradient descent approaches: compute the gradient of the loss w.r.t. the parameters **\mathbf{w}** and apply gradient descent as usual
 - For example: SciKit Learn **`linear_model.SGDClassifier()`** with parameters: `loss='hinge'` and `penalty='l2'`. Note: instead of parameter C giving weight to the classification performance, they use parameter alpha giving weight to the regularisation term.
 - <http://scikit-learn.org/stable/modules/sgd.html#sgd>
 - http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
 - Practical considerations: we need to tune the learning rate, could be advantageous for very large training sets and cases where incremental learning is needed

7.3 Nonlinear problems


Nonlinear problems

- **Problem:** how to linearly separate these 2 classes of samples?
- 2 solutions
 - move to non-linear decision boundaries by adding non-linear features to the x array and then use a linear SVM as usual - see slide 7
 - use non-linear SVMs with **kernels**
 - **see next slides**




Nonlinear SVM for not linearly separable data

- **Map** the sample **input space** to a higher dimensional space (called **feature space**) with a function ϕ ...



Nonlinear SVM for not linearly separable data

- ...where samples are linearly separable by SVM !



Nonlinear SVM for not linearly separable data

- In summary, the SVM problem remains the same as before except that \mathbf{x} is replaced by $\phi(\mathbf{x})$ in all equations
- The vector \mathbf{w} defining the **optimal hyperplane** is found through the learning process using $\mathbf{w} \cdot \phi(\mathbf{x}) + b = 0$
- Then a new sample \mathbf{x}_t is classified by computing $\text{sign}(\mathbf{w} \cdot \phi(\mathbf{x}_t) + b)$

In red what changes from previous methods


Nonlinear SVM for not linearly separable data

- The functions ϕ are computed through *kernel functions* located at each training points \mathbf{x}_i
 - A new test sample \mathbf{x}_t is classified by computing the sign of $\mathbf{w} \cdot \phi(\mathbf{x}_t) + b = \sum \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}_t) + b$
- **Common kernels** $K(\mathbf{x}_i, \mathbf{x}_j)$: ($i, j = 1, \dots, N$)
 - Linear: $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i \cdot \mathbf{x}_j$
 - Polynomial: $K(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i \cdot \mathbf{x}_j + 1)^d$ (where d is the polynom degree)
 - Gaussian or radial basis function (RBF):
 $K(\mathbf{x}_i, \mathbf{x}_j) = \exp(-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|^2)$ (where $\gamma = 1/2\sigma^2 > 0$)
 - Hyperbolic tangent:
 $K(\mathbf{x}_i, \mathbf{x}_j) = \tanh(k \mathbf{x}_i \cdot \mathbf{x}_j + c)$ (for some $k > 0$ and $c < 0$)

7.4 Multiclass SVM


Multiclass SVM

- Initially, **SVM** are a **binary classifier**, i.e. can separate 2 classes of samples
- Problem:** how to separate multiple classes?




Multiclass SVM

- **Solution:** reduce the single multiclass problem into **multiple binary classification problems**




Multiclass SVM

- **One-vs-all method:** the classification of new samples is done by a winner-takes-all strategy, in which **the SVM with the highest output value assigns the class to a given sample**
- For a given test sample \mathbf{x}_t , the output value is given by the function $\mathbf{w}_c \cdot \mathbf{x}_t + b_c$, where $c = 1, \dots, M$ and M is the number of classes



Multiclass SVM

- **One-vs-one method:** the classification of new samples is done by a max-wins voting strategy, in which **every SVM classifier assigns a given sample to one of the two classes**, then the number of votes for the assigned class is increased by one, and finally **the class with the highest number of votes is assigned to the sample**



7.5 History & References

History of SVM

- **1960**: Beginning of SVM development
- **1963**: Original **linear SVM** algorithm proposed by Vladimir N. Vapnik
- **1964**: **Kernel trick** first published by M. Aizerman, E. Braverman, and L. Rozonoer
- **1992**: **Nonlinear SVM** (using **kernel trick**) proposed by Bernhard E. Boser, Isabelle M. Guyon and Vladimir N. Vapnik
- **1995**: Current **standard SVM** (using **soft margin**) proposed by Vladimir N. Vapnik and Corinna Cortes
- **1996**: SVM for regression (**SVR**) proposed by Vladimir N. Vapnik
- **2008**: Vladimir N. Vapnik and Corinna Cortes received the ACM Paris Kanellakis Award for their scientific contribution
- **Today**: SVM are widely used because of their efficiency

References

- **SVM on Wikipedia:**

http://en.wikipedia.org/wiki/Support_vector_machine

- **SVM tutorials:**

<http://www.svms.org/tutorials>

<http://www.kernel-machines.org/tutorials>

- **Nice SVM presentation (with biomedical application):**

A. Statnikov, D. Hardin, I. Guyon, C. F. Aliferis,

A Gentle Introduction to Support Vector Machines in Biomedicine,

<http://www.med.nyu.edu/chibi/sites/default/files/chibi/Final.pdf>

- **Books:**

C. Bishop, *Pattern Recognition and Machine Learning*, Springer, 2006

R. O. Duda, P. E. Hart, D. G. Stork, *Pattern Classification (2nd Ed.)*,

Wiley-Interscience, 2001

